

TUGRUL GUNER

AI Engineering Leader | Portfolio Strategy | Hands-on Architecture

Toronto, ON, Canada | tugrulgunr@gmail.com | (514) 585-3990 | linkedin.com/in/tugrulguner

PROFILE

AI engineering leader and hands-on architect who builds high-performing teams and scalable, robust backend, AI, and agentic systems. Lead two cross-functional pods across a portfolio of more than six products, combining product direction and stakeholder leadership with architecture and production delivery. Operate well under pressure and make difficult tradeoffs without lowering engineering standards.

EXPERIENCE

AI Engineering Manager | CDAI

Toronto, Ontario, Canada | Jan 2026–present

- Lead two separate cross-functional pods of five and seven people across a portfolio of more than six products, with responsibility for technical direction, cross-team alignment, and delivery.
- In under six months, led one pod to a 2.5× increase in delivery velocity and helped the team launch three products. That delivery improvement led to responsibility for the second, separate pod.
- Helped deliver a recommendation system that improved menu-item attachment rate by 4×.
- Remain hands-on: prototype ideas, make architectural changes, deliver production work, and build shared platforms that reduce redundancy and improve scalability.
- Build high-performing teams around ownership and engineering standards; started the AI & Science channel and co-host the company-wide AI Center of Excellence meeting.

Technical Manager, AI Architect | Haptiq

Toronto, Ontario, Canada | May 2025–Jan 2026

- Built Haptiq's AI team and first AI infrastructure from the ground up; recruited and led seven people across ML engineering, MLOps, and backend development.
- Led technical delivery and client conversations for the company's first AI product, helping sell AI solutions to companies and contributing to more than \$2M in new revenue.
- Worked with clients and executives on product direction, architecture, delivery plans, risk assessments, and performance reporting.
- Established ownership, knowledge-sharing, reviews, documentation, release management, and sprint planning for reliable delivery.

Senior Tech Lead | Quotograph.io

Waterloo, Ontario, Canada · Remote | Jul 2025 – Present

Chief Technology Officer | Quotograph.io

Waterloo, Ontario, Canada · Remote | Jan 2024 – Jul 2025

Lead MLOps Engineer | Arteria AI

Toronto, Ontario, Canada | Dec 2024–May 2025

- Architected scalable AI and backend systems aligned with business goals and integrated Ray with Kubernetes.
- Led cross-functional delivery, mentoring, cost optimization, availability, and fault-tolerance work.

Senior MLOps Engineer | Arteria AI

Toronto, Ontario, Canada | Sep 2022–Dec 2024

- Automated backend systems for model integration and scaling, and built user-friendly APIs for rapid experimentation and deployment.
- Reduced Kubernetes infrastructure costs by 60% with KServe and improved model performance 10× with quantization and ONNX.
- Built CI/CD pipelines with GitHub Actions and Databricks to streamline the machine-learning lifecycle.

TECHNICAL RANGE

Languages: Python, SQL, Bash, Git | **ML:** PyTorch, TensorFlow, Hugging Face, DSPy, vLLM

Agents: Hermes Agent, Superpowers, Deep Agents, Pydantic AI, LangGraph, MCP/FastMCP | **MLOps:** KServe, MLflow, Ray Serve, LitServe

Platforms: AWS, Azure, GCP, Docker, Kubernetes, Databricks, GitHub Actions | **Systems:** FastAPI, Django, PostgreSQL, Qdrant, MongoDB, Redis, RabbitMQ

EDUCATION & RESEARCH

Postdoctoral Fellow, University of Ottawa, 2021–2022 — deep learning for quantum optics and beam shaping

Postdoctoral Fellow, Institut national de la recherche scientifique (INRS), 2018–2021 — ultrafast transmission electron microscopy and measurement algorithms

MBA, Quantic School of Business and Technology, 2025 | **Ph.D., Materials Science and Engineering**, Izmir Institute of Technology, 2018

30 journal publications, one patent, 1,023 Google Scholar citations, h-index 19, i10-index 21 (September 2026).

OPEN SOURCE & WRITING

Created **ModePot**, an open-source family focused on simpler APIs, modern runtimes, performance, and bounded agent-owned decisions. Contribute to **Hermes Agent** and **Superpowers**. Write **Passionately Curious** and publish essays on Medium.