

Tugrul Guner

tugrulgunr@gmail.com | (514) 585-3990 | Toronto, ON, Canada | [linkedin.com/in/tugrulguner](https://www.linkedin.com/in/tugrulguner)

Professional Summary

Strategic AI and MLOps leader with an MBA and over five years of experience architecting and scaling AI/ML systems, supported by a decade of Python expertise. Currently serving as AI Engineering Manager at CDAI. Responsible for setting the AI roadmap, designing enterprise-grade architectures, building and managing a cross-functional engineering team, and aligning initiatives with executive and business priorities.

Proven track record of translating business strategy into scalable, cost-efficient AI solutions, driving measurable revenue impact, and establishing engineering practices that accelerate delivery while reducing technical debt. Combining technical leadership with business strategy, I am positioned to scale AI capabilities and lead organizational transformation.

Technical Skills

- **Programming Languages:** Python, SQL, Bash, Git
- **Machine Learning Frameworks:** TensorFlow, PyTorch, HuggingFace
- **Agent Frameworks:** pydantic-ai, LangGraph
- **LLM Frameworks:** LlamaIndex, DSPy, vLLM, crewAI
- **MLOps Tools:** KServe, MLflow, Ray Serve, LitServe
- **Cloud Platforms:** AWS, Azure, GCP
- **DevOps Tools:** Docker, Kubernetes, GitHub Actions, Azure DevOps, Datadog
- **Databases:** PostgreSQL, Qdrant, MongoDB, ChromaDB
- **Others:** Databricks, Celery, RabbitMQ, Redis, FastAPI, FastMCP, Django, Ray

Selected Work Experience

AI Engineering Manager, CDAI

Toronto, ON, Canada (Hybrid, Full-time)

01/2026 - Present

- **Lead a cross-functional pod** of 8 engineers, designers, and product stakeholders, owning delivery end-to-end
- **Drive technical direction and modernization** of a production system serving thousands of active users
- **Shape product roadmap and translate business goals** into scalable technical initiatives
- **Hire, mentor, and develop** team members to improve execution velocity and engineering quality

Technical Manager, AI Architect, Haptiq

Toronto, ON, Canada (Hybrid, Full-time)

05/2025 - 01/2026

- **Founded and built Haptiq's AI division**, creating the company's AI roadmap, architecture, and engineering practices from the ground up.
- **Delivered \$2M+ in new revenue** by designing and deploying agentic AI capabilities through modular AI services and external MCP integrations.

- **Recruited, onboarded, and led** a cross-functional team of ML engineers, MLOps specialists, and backend developers with team size of 7; established a culture of collaboration, knowledge-sharing, and accountability.
- **Architected scalable, modular, and fault-tolerant AI infrastructure**, enabling rapid experimentation and reliable production deployment.
- **Standardized engineering practices**, including code reviews, documentation, and release management, reducing technical debt and accelerating feature delivery.
- **Drove project execution**, leading sprint planning, milestone setting, and clear ownership to ensure on-time delivery of AI-powered initiatives.
- **Provided executive-level reporting**, presenting project status, risk assessments, and performance metrics (model accuracy, deployment latency, infrastructure health) to C-suite leadership.

Lead MLOps Engineer, Arteria AI

Toronto, ON, Canada (Remote, Full-time)

12/2024 - 05/2025

- **Architect and implement** scalable AI and backend solutions aligned with business goals.
- **Mentor and guide** cross-functional teams for seamless collaboration.
- **Optimize costs**, reducing infrastructure expenses and improving ROI.
- **Integrate Ray with Kubernetes**, enabling scalable, cost-effective AI deployments.
- **Promote automation and best practices**, fostering continuous improvement.
- **Ensure high availability and fault tolerance** in AI deployments.

Senior MLOps Engineer, Arteria AI

Toronto, ON, Canada (Remote, Full-time)

09/2022 - 12/2024

- **Automated 100% of backend systems**, accelerating AI model integration and scaling.
- **Developed user-friendly APIs**, enabling rapid model experimentation and deployment.
- **Optimized Kubernetes-based AI deployments**, cutting infrastructure costs by 60% with KServe.
- **Enhanced model performance 10x** using quantization and ONNX optimizations.
- **Built CI/CD pipelines** with GitHub Actions and Databricks, streamlining ML lifecycle management.

Education

MBA, Quantic School of Business and Technology

08/2025

Ph.D., Materials Science and Engineering, Izmir Institute of Technology

12/2018

Publications & Patents

- **30 Publications** in International Scientific Journals including Nature, Elsevier, IOP, etc. and 1 **Patent**
- **>900 Citations**, **h-index: 18** and **i-10: 20**