

TUGRUL GUNER

AI Engineering Leader | Portfolio Strategy | Hands-on Architecture

Toronto, ON, Canada | tugrulgunr@gmail.com | (514) 585-3990 | linkedin.com/in/tugrulguner

PROFILE

AI engineering leader and hands-on architect who builds high-performing teams and scalable, robust backend, AI, and agentic systems. Lead two cross-functional pods across a portfolio of more than six products, combining product direction and stakeholder leadership with architecture and production delivery. Operate well under pressure and make difficult tradeoffs without lowering engineering standards.

EXPERIENCE

AI Engineering Manager | CDAI

Toronto, ON | Jan 2026 - Present

- Lead two separate cross-functional pods of five and seven across more than six products, owning technical direction, alignment, and delivery.
- In under six months, one pod achieved 2.5× delivery velocity and the team launched three products; the result led to responsibility for a second, separate pod.
- Helped deliver a recommendation system that improved menu-item attachment rate 4×.
- Remain hands-on across architecture and production work; co-host the company-wide AI Center of Excellence meeting.

Technical Manager, AI Architect | Haptiq

Toronto, ON | May 2025 - Jan 2026

- Built the AI team and first AI infrastructure, recruiting and leading seven ML, MLOps, and backend engineers.
- Led technical delivery and client conversations for the first AI product, contributing to more than \$2M in new revenue.
- Shaped product direction, architecture, delivery plans, and engineering practices with clients and executives.

Senior Tech Lead | Quotograph.io

Waterloo, Ontario, Canada (Remote) | Jul 2025 - Present

Chief Technology Officer | Quotograph.io

Waterloo, Ontario, Canada (Remote) | Jan 2024 - Jul 2025

Lead MLOps Engineer | Arteria AI

Toronto, ON | Dec 2024 - May 2025

- Architected scalable AI and backend systems and integrated Ray with Kubernetes.
- Led cross-functional delivery, mentoring, cost optimization, availability, and fault tolerance.

Senior MLOps Engineer | Arteria AI

Toronto, ON | Sep 2022 - Dec 2024

- Automated model integration/scaling and built APIs for experimentation and deployment.
- Cut Kubernetes costs 60% with KServe; improved model performance 10× with quantization and ONNX.
- Built ML lifecycle CI/CD with GitHub Actions and Databricks.

TECHNICAL RANGE

Languages: Python, SQL, Bash, Git | **ML:** PyTorch, TensorFlow, Hugging Face, DSPy, vLLM

Agents: Hermes Agent, Superpowers, Deep Agents, Pydantic AI, LangGraph, MCP/FastMCP | **MLOps:** KServe, MLflow, Ray Serve, LitServe

Platforms: AWS, Azure, GCP, Docker, Kubernetes, Databricks, GitHub Actions | **Systems:** FastAPI, Django, PostgreSQL, Qdrant, MongoDB, Redis, RabbitMQ

EDUCATION & RESEARCH

Postdoctoral Fellow, University of Ottawa, 2021–2022 — deep learning for quantum optics and beam shaping

Postdoctoral Fellow, Institut national de la recherche scientifique (INRS), 2018–2021 — ultrafast transmission electron microscopy and measurement algorithms

MBA, Quantic School of Business and Technology, 2025 | **Ph.D., Materials Science and Engineering**, Izmir Institute of Technology, 2018

30 journal publications, one patent, 1,023 Google Scholar citations, h-index 19, i10-index 21 (September 2026).

OPEN SOURCE & WRITING

Created **ModePot**, an open-source family focused on simpler APIs, modern runtimes, performance, and bounded agent-owned decisions. Contribute to **Hermes Agent** and **Superpowers**. Write **Passionately Curious** and publish essays on Medium.